

ElderTech Voice Assistant – Full-Stack App Blueprint

Purpose

Create a warm, **voice-first** assistant that helps older adults master everyday technology through natural conversation and spoken answers—no animated avatar required.

1. High-Level User Flow

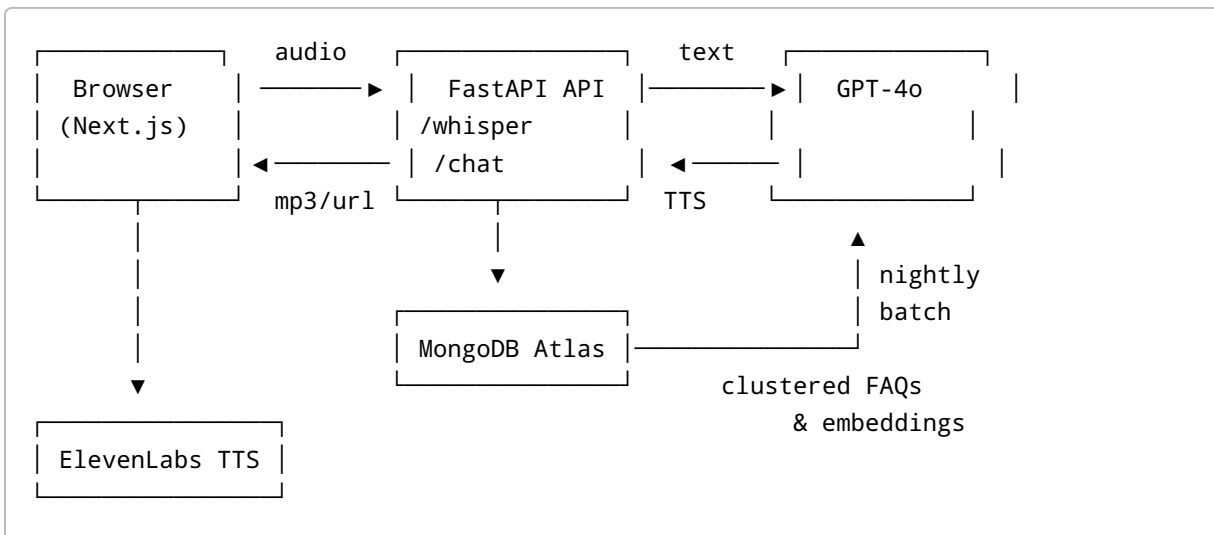
- 1. **Launch:** User opens the web app on a tablet/desktop. First-time visitors see an onboarding walkthrough with large visuals and voice narration.
- 2. **Conversation:** A friendly assistant icon gently pulses, prompting the user to tap a large **“Hold to Talk”** button (or use a wake-word).
- 3. **Transcription:** Microphone audio streams to the backend via WebRTC; Whisper API transcribes in real-time.
- 4. **NLP Processing:** Backend sends conversation context + transcript to GPT-4o with an empathy-focused system prompt tuned for seniors.
- 5. **Response Rendering:**
- 6. GPT returns plain text/markdown.
- 7. FastAPI calls **ElevenLabs TTS** → returns an MP3 of the spoken answer.
- 8. **Playback:** Front-end shows 24-pt subtitles and plays the audio. Controls: **Replay | Slower | Louder**.
- 9. **Logging & Learning:** Q&A pair saved to MongoDB → nightly batch clusters similar questions, grows FAQs and RAG index.

2. Tech Stack

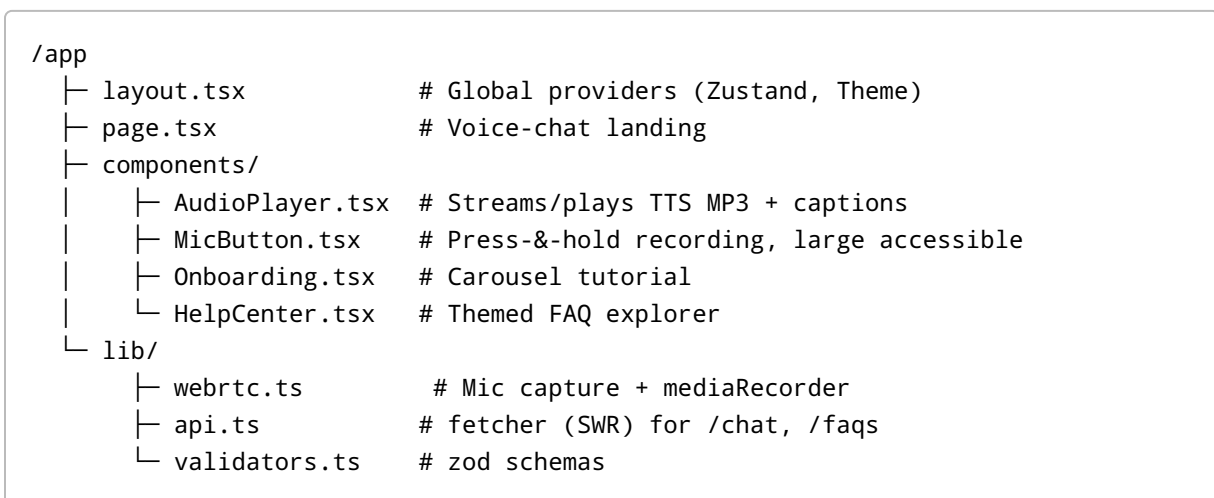
Layer	Choice	Rationale
Frontend	Next.js 14 (React Server Components) Tailwind CSS Zustand	Fast, accessible UI with clear state management.
Real-time	WebRTC (audio only) Socket.IO fallback	Low-latency mic streaming; works on restrictive networks.
Backend API	FastAPI + Pydantic v2	Async, type-safe, automatic OpenAPI docs.
Database	MongoDB Atlas (<code>users</code> , <code>sessions</code> , <code>messages</code> , <code>faqs</code>)	Flexible document model fits conversational data.

Layer	Choice	Rationale
AI Services	OpenAI Whisper & GPT-4o; ElevenLabs TTS	Best-in-class transcription, reasoning, and natural-sounding speech.
DevOps	Docker Compose (dev) → Render.com / Fly.io / Vercel Edge (prod)	Simple local spin-up; global edge functions for low latency.

3. System Architecture



4. Frontend Breakdown (Next.js `/app` directory)



Key UI Patterns

- **Large Typography:** 18-24 pt base font; 44 px touch targets.
- **Voice Everywhere:** `aria-label` + optional spoken hints on all controls.
- **Assistive Controls:** *Repeat, Slower, Examples* buttons always visible.

5. Backend Breakdown (FastAPI)

```
backend/
├─ main.py           # FastAPI app factory
├─ routers/
│   ├─ chat.py       # /chat POST {transcript}
│   ├─ whisper.py     # /whisper WS stream
│   ├─ faqs.py        # CRUD on FAQs
│   └─ auth.py        # JWT /family login
├─ services/
│   ├─ openai.py      # GPT-4 & Whisper calls
│   ├─ tts.py         # ElevenLabs wrapper
│   └─ db.py          # Motor (async Mongo) client
└─ scripts/
    └─ nightly_faq_clustering.py
```

Sample `/chat` Endpoint

```
@router.post("/chat", response_model=ChatResponse)
async def chat(req: ChatRequest, user: User = Depends(require_user)):
    # 1. Generate answer
    system = "You are a patient tech coach..."
    msgs = [
        {"role": "system", "content": system},
        *req.history,
        {"role": "user", "content": req.transcript},
    ]
    gpt_resp = await openai.chat(messages=msgs, model="gpt-4o")
    answer = gpt_resp.choices[0].message.content

    # 2. Generate speech
    mp3_url = await tts.speak(answer)

    # 3. Persist
    await db.messages.insert_one({
        "user_id": user.id,
        "q": req.transcript,
        "a": answer,
```

```
        "audio": mp3_url,  
        "ts": datetime.utcnow()  
    })  
  
    return {"answer": answer, "audio": mp3_url}
```

6. Database Schema (Mongo)

```
users: {  
  _id, name, role: "elder" | "family", language, createdAt  
}  
sessions: {  
  _id, user_id, startedAt, endedAt  
}  
messages: {  
  _id, session_id, q, a, audio, ts  
}  
faqs: {  
  _id, question, answer_md, tags: ["email", "scams"], updatedAt  
}
```

7. Accessibility & UX Principles

- **WCAG 2.2 AA** compliant colors, text sizes, captions.
- **Keyboard & switch-control** navigation (focus rings, `tabindex`).
- **Latency budget:** <400 ms TTFB, <2 s TTS playback (cached path).

8. Onboarding Experience

1. Welcome screen with assistant voice: “Hi! I’m here to help you with tech.”
2. Guided demo: User taps mic, asks “What’s an app?” → hears sample reply.
3. Brief quiz: User tries asking; gets celebratory sound + confetti.
4. Option to browse “common topics” grid or start chatting.

9. Family Portal (Optional)

- Auth0-protected dashboard.
- Pre-load FAQs, view transcripts & audio, flag incorrect answers.
- Email digests of weekly usage.

10. Deployment & Ops

- **Dev:** `docker compose up` spins Next.js + FastAPI + Mongo.
 - **Staging:** Vercel (frontend) → Render (backend) → MongoDB Atlas.
 - **Observability:** Sentry (front & back), Prometheus + Grafana metrics.
 - **CI/CD:** GitHub Actions → lint, type-check, Playwright e2e.
-

11. Next Steps

1. Scaffold repo with `pnpm create next-app@latest` + `fastapi-project-generator`.
2. Implement Whisper streaming endpoint.
3. Build **AudioPlayer** component and connect to `/chat`.
4. Conduct hallway tests with 2-3 seniors; refine controls & font sizes.
5. Seed FAQ JSON; connect nightly clustering job later.

-- End of Blueprint --